

Machine Learning

Regression, Classification

| Context: How Machines Actually Learn

The Optimization Loop

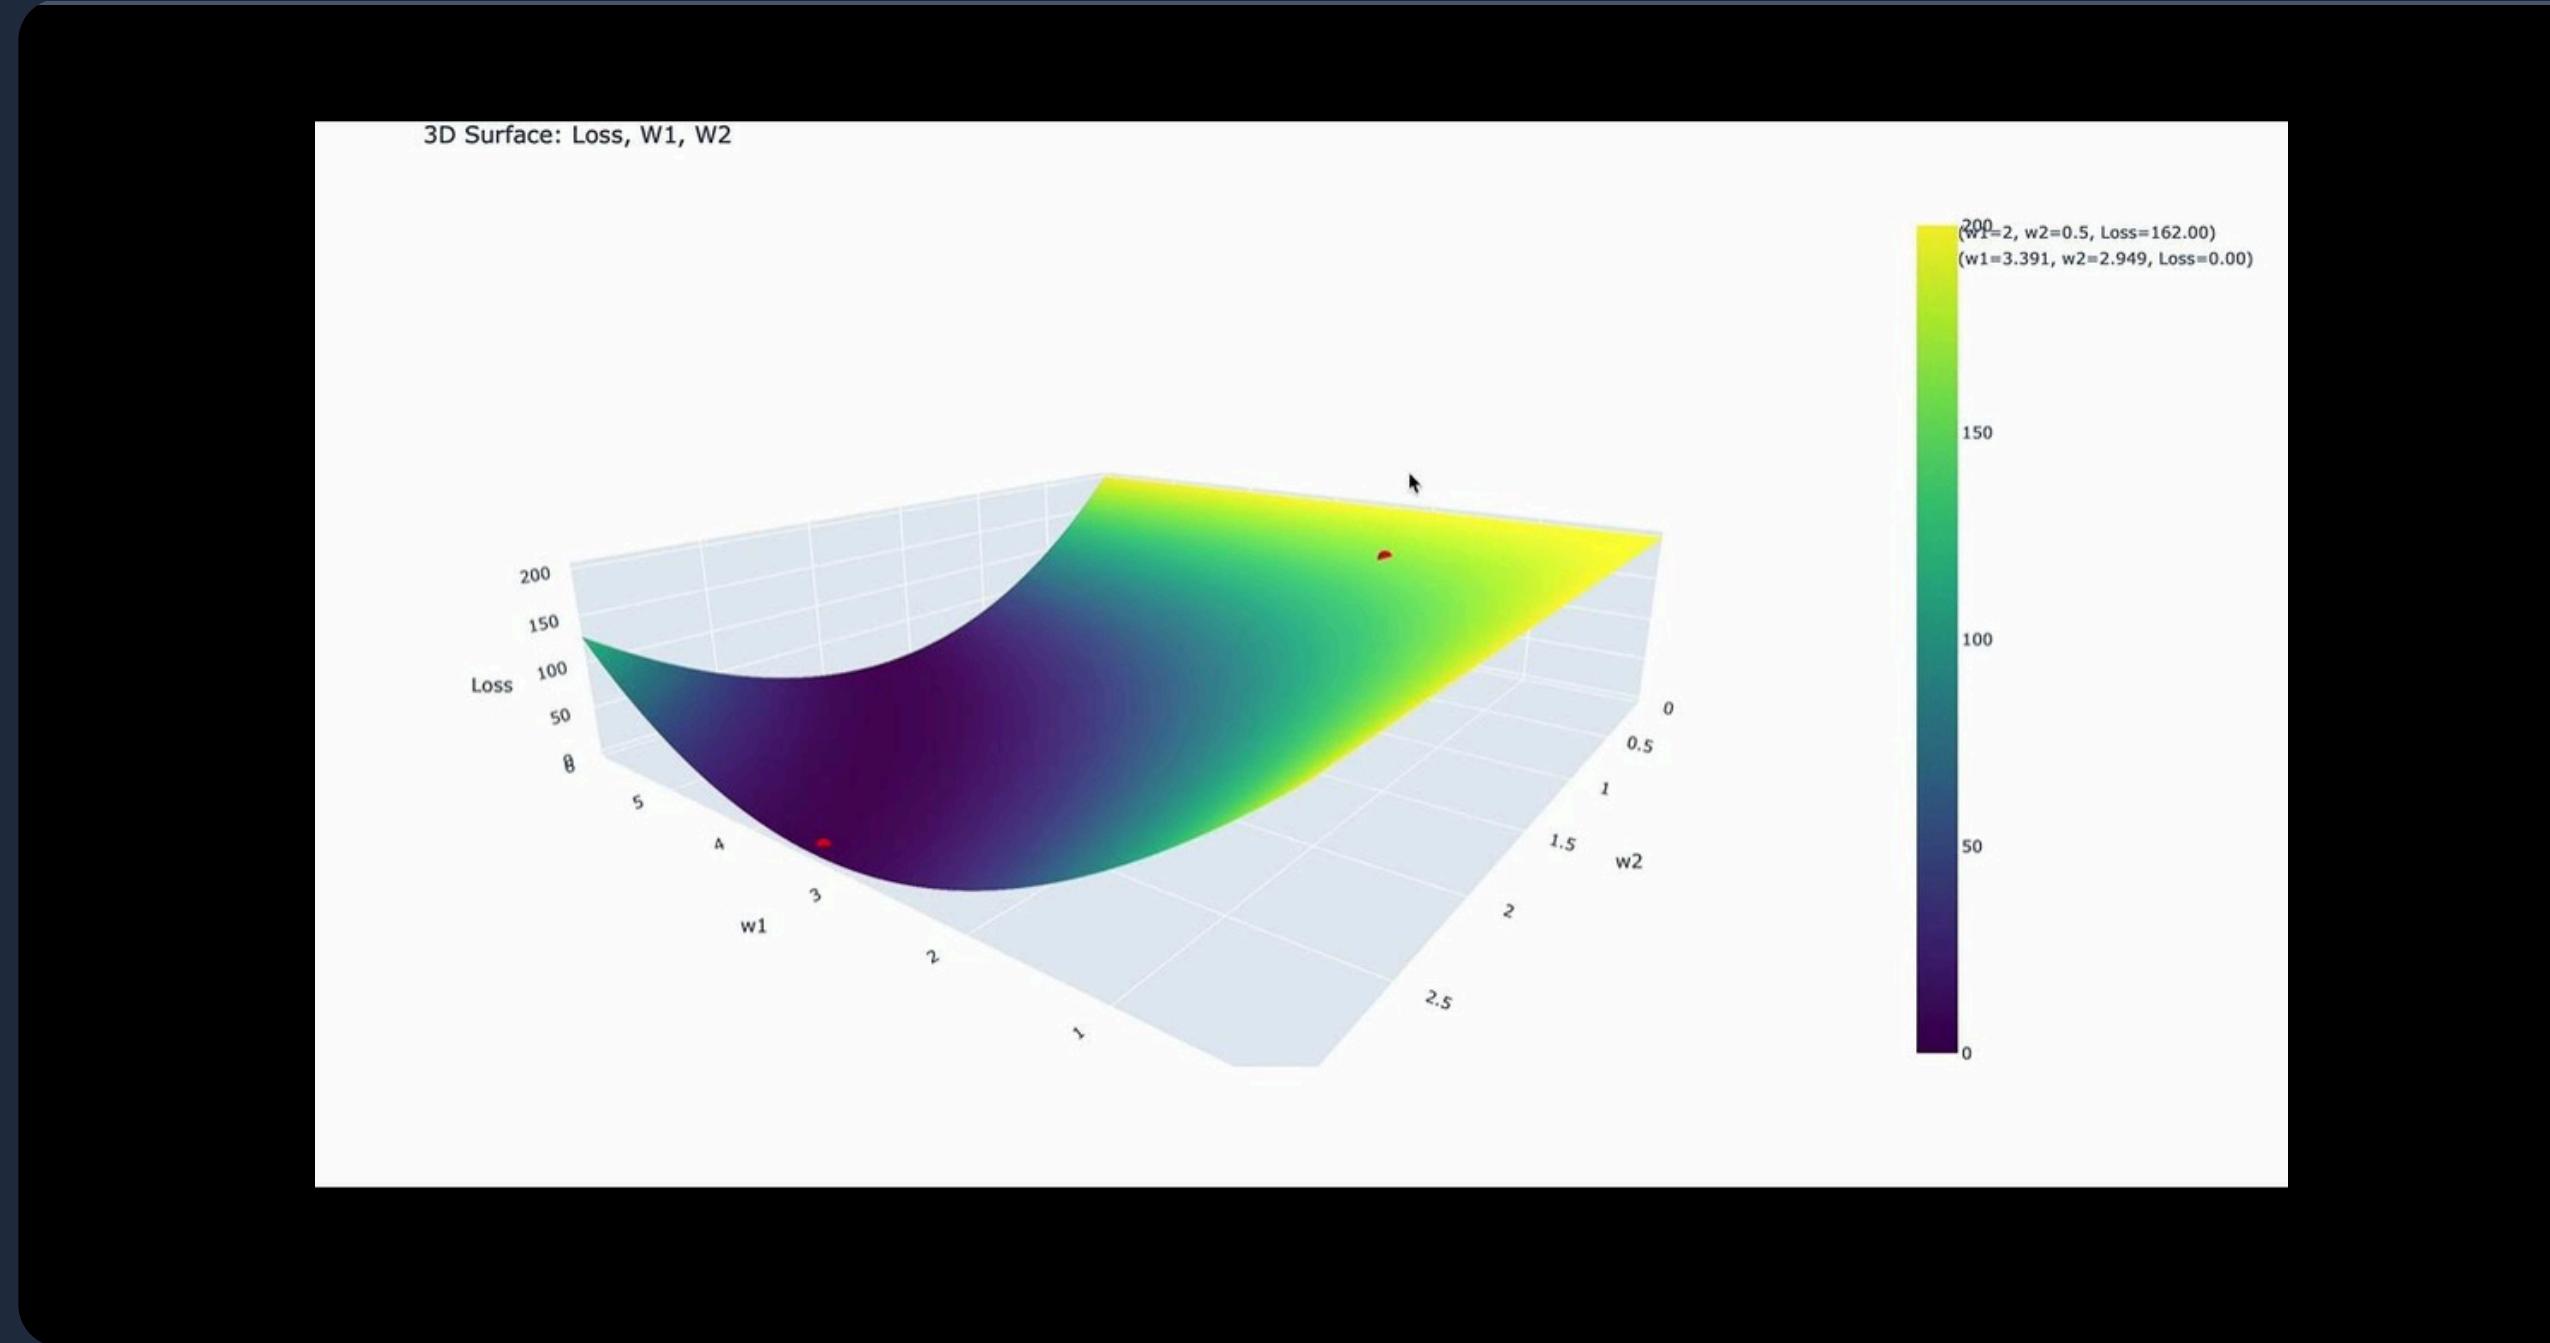
Machine learning isn't magic; it's math. The goal is to minimize error.

- **1. Forward Pass:** Model makes a guess.
- **2. Loss Function:** Calculates how wrong the guess was (Error).
- **3. Optimizer:** Adjusts the weights to reduce error (e.g., Gradient Descent).
- **4. Repeat:** Until the error stops decreasing.

Key Concept: Weights & Biases

The "brain" of the model is just a list of numbers (parameters). "Learning" simply means finding the best numbers for that list.

Visualizing Optimization



Gradient Descent: The algorithm walks "downhill" on the error surface to find the global minimum (lowest error).

Step 0: Preprocessing

Train / Test Split

We hide data from the model to test it later. Usually an 80/20 split.

```
{ from sklearn.model_selection import train_test_split  
x_train, x_test = train_test_split(x, test_size=0.2)
```

Feature Scaling

Algorithms like KNN and SVM break if numbers are on different scales (e.g., Age vs Salary).

```
{ from sklearn.preprocessing import StandardScaler  
scaler = StandardScaler()  
X_scaled = scaler.fit_transform(X)
```

| The Two Main Paradigms

Supervised

Labeled Data (Inputs + Answers)

We teach the model what to predict.

- Regression (Price)
- Classification (Spam/Not Spam)

Unsupervised

Unlabeled Data (Inputs Only)

The model finds hidden patterns on its own.

- Clustering (Customer segments)
- Dimensionality Reduction

| Topic 1: Regression

Predicting Quantity

Regression maps input variables to a **continuous** output.

- Stock Prices
- Temperature
- House Value

$$y = f(x) + \varepsilon$$

Linear Regression

The Equation

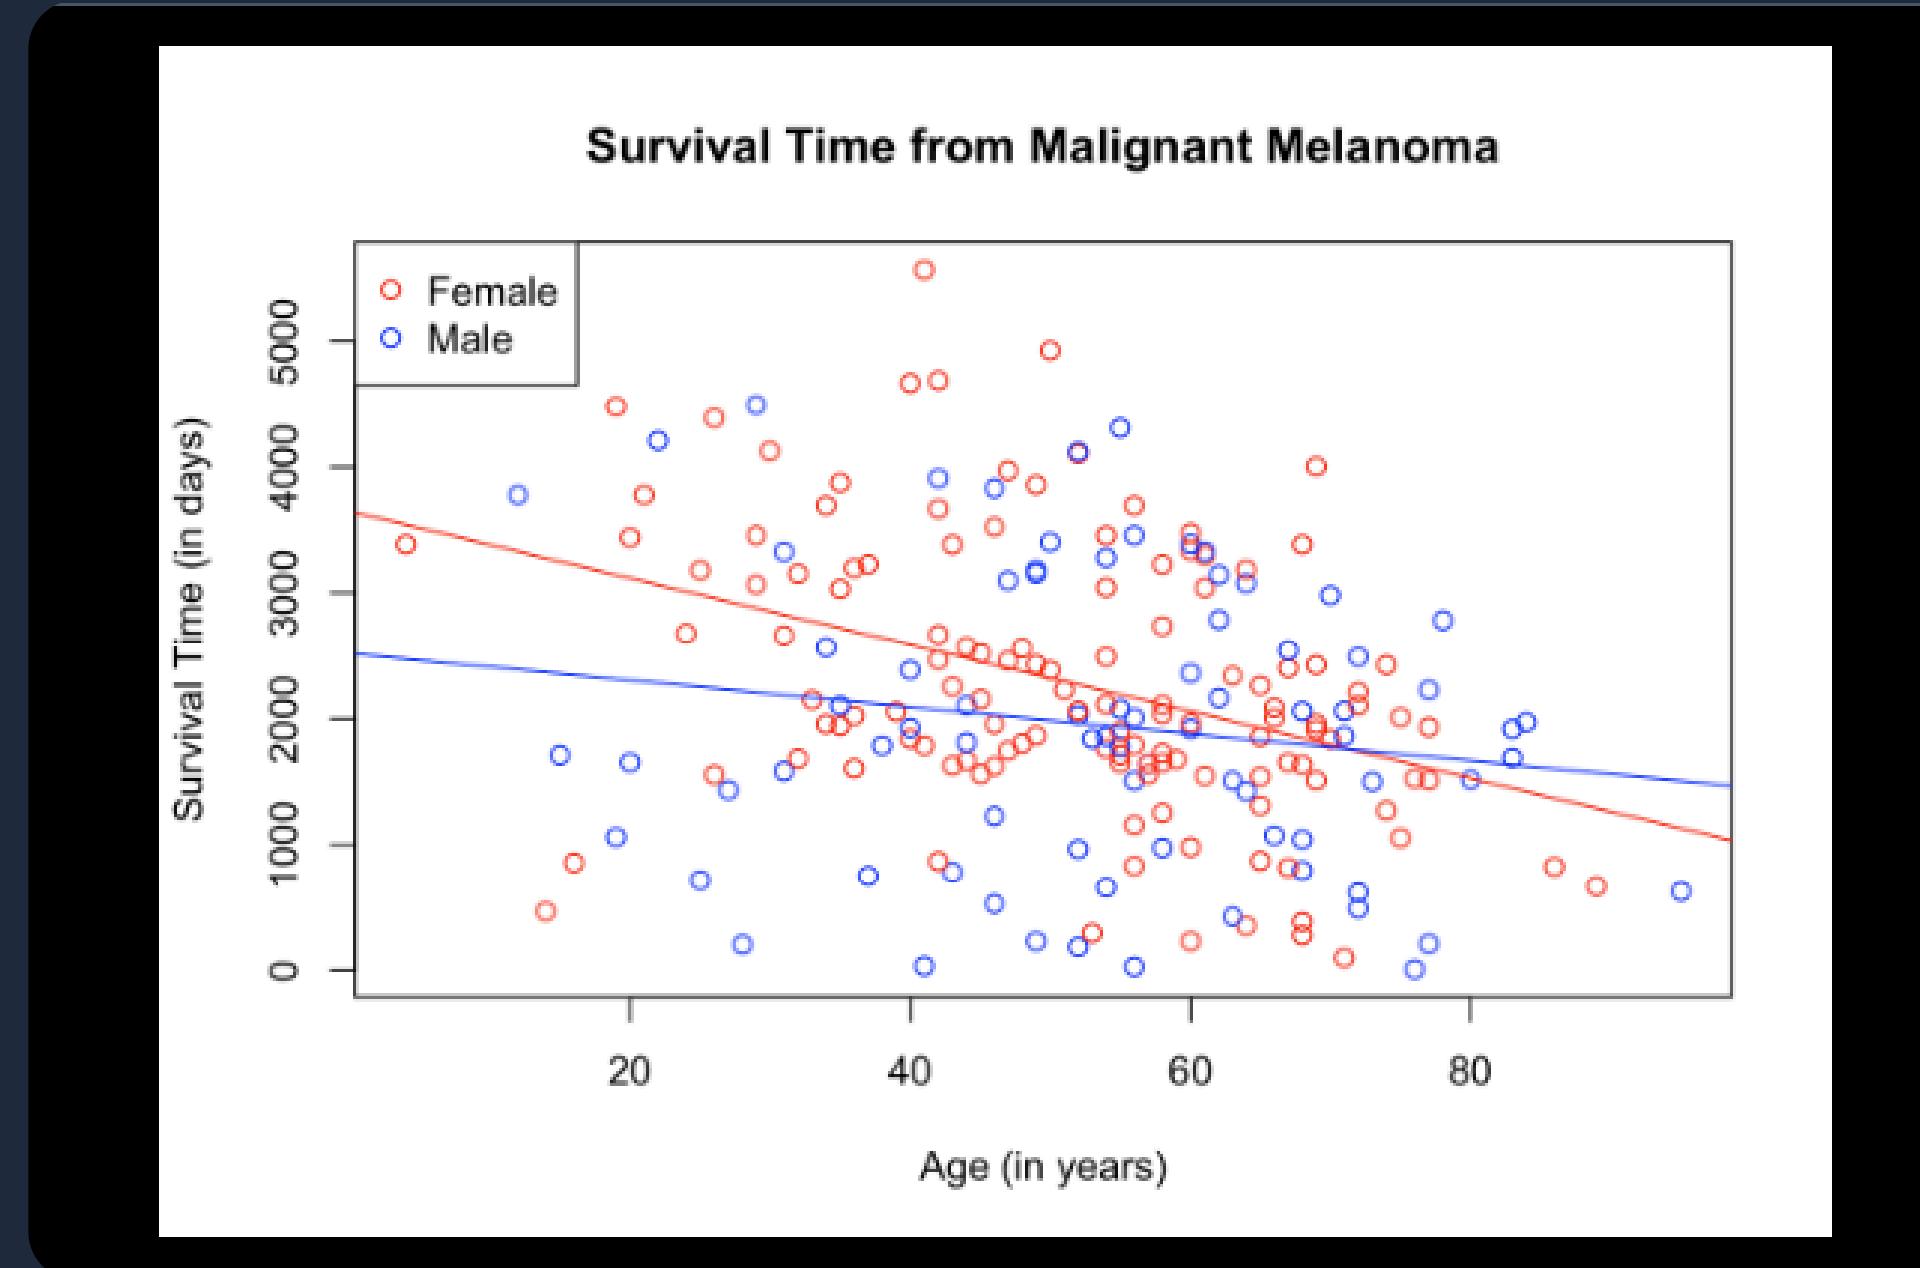
Fitting a straight line through data points.

$$y = \beta_0 + \beta_1 x$$

The model learns (β_0) (intercept) and (β_1) (slope) to minimize the distance between the line and the dots.

```
from sklearn.linear_model import LinearRegression
model = LinearRegression()
model.fit(X_train, y_train)
# Predicts a continuous value
preds = model.predict(X_test)
```

Visualizing the Fit



The blue line represents the model's predictions. The vertical distance from each dot to the line is the **Residual (Error)**.

| Non-Linear Regression

Polynomial Features

Not everything is a straight line. By adding powers of X (x^2 x^3), we can model curves.

Warning: Higher degrees make the model more complex, risking *Overfitting*.

```
from sklearn.preprocessing import PolynomialFeatures  
from sklearn.pipeline import make_pipeline # Creates x, x^2,  
x^3 features  
poly = make_pipeline(PolynomialFeatures(3),  
LinearRegression())  
poly.fit(X_train, y_train)
```

| Context: Bias-Variance Tradeoff

Underfitting (High Bias)

Model is too simple. It misses the trend.

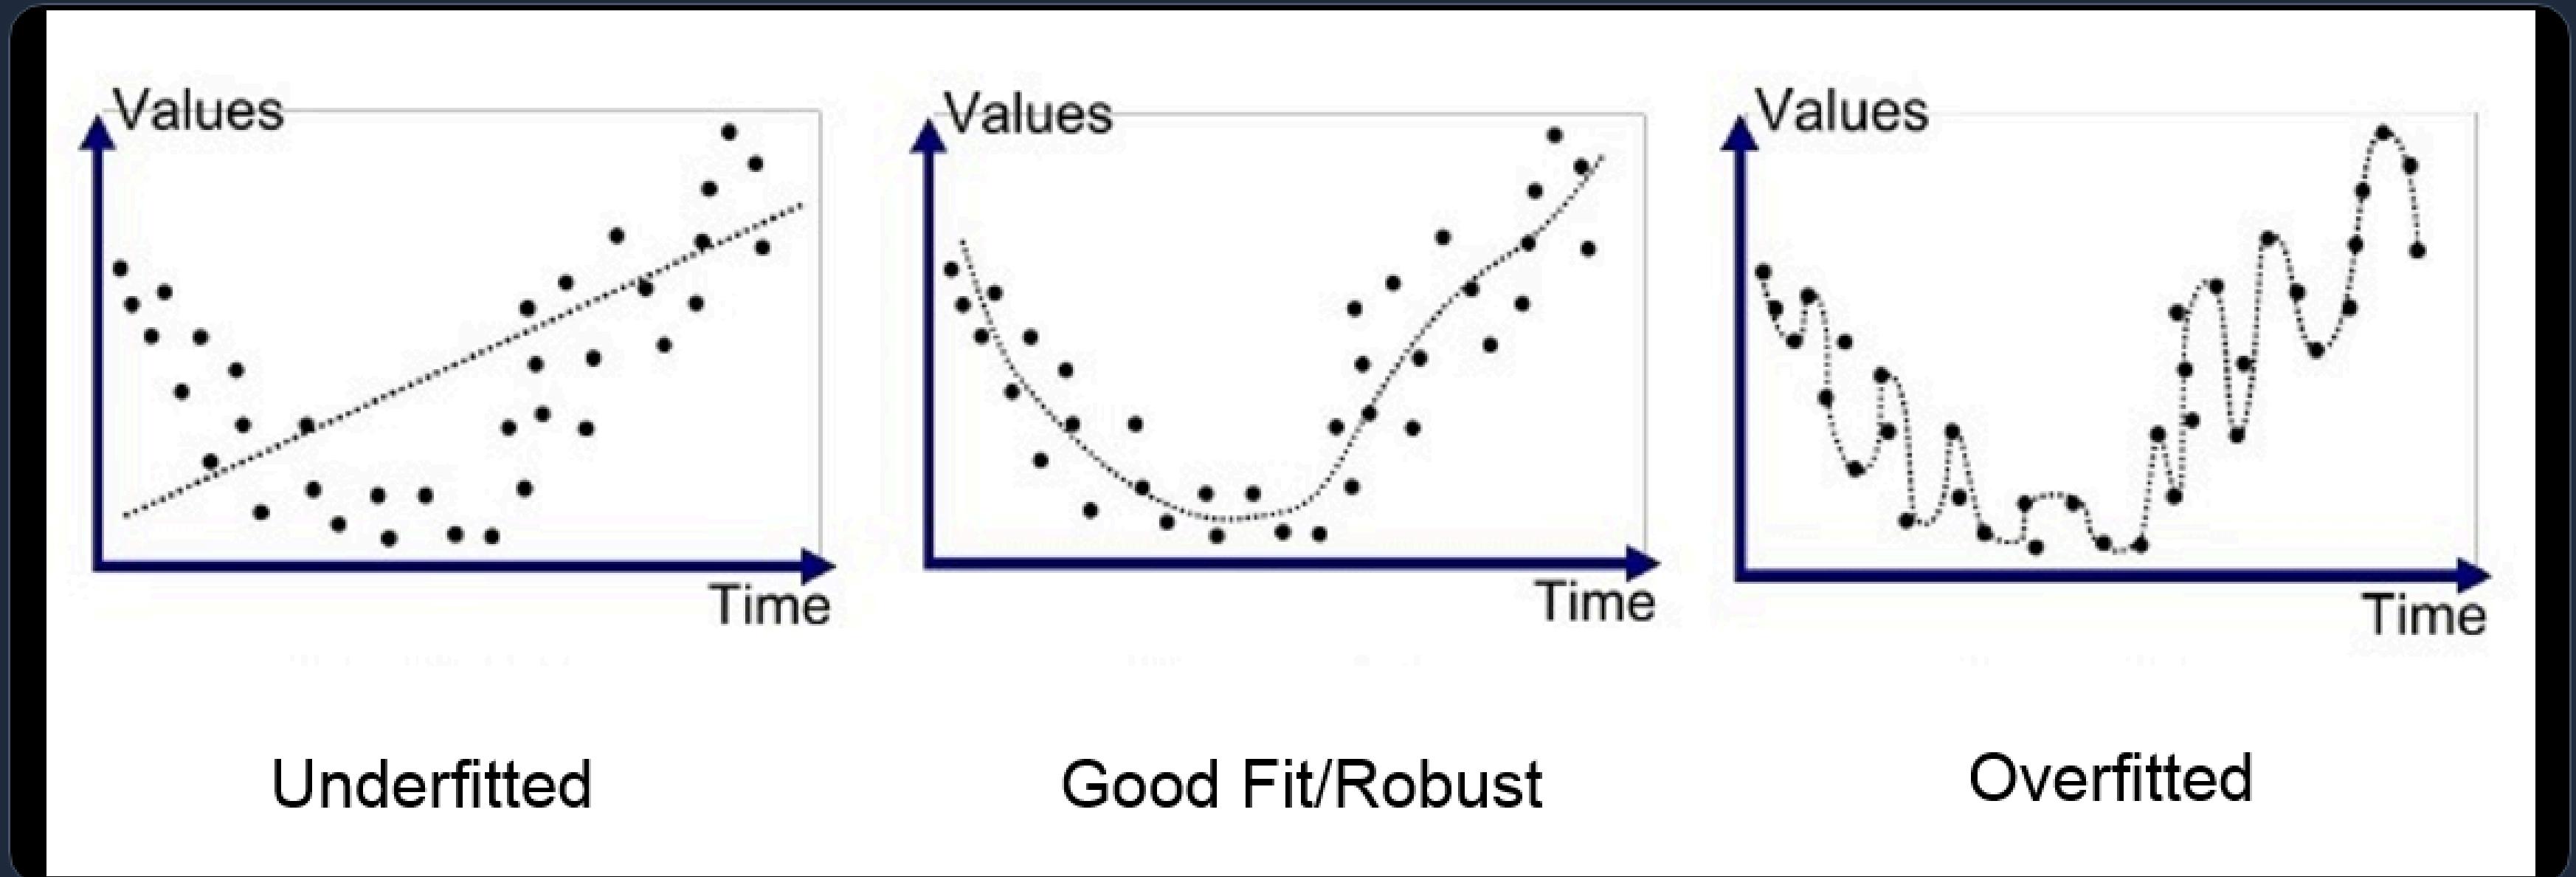
Example: Fitting a straight line to a parabola.

Overfitting (High Variance)

Model is too complex. It memorizes noise.

Example: A squiggly line that touches every single dot.

The Sweet Spot



Our goal is the middle graph: Capturing the general trend without capturing the random noise.

Metrics: MAE vs MSE

MAE (Mean Absolute Error)

Average of absolute errors.

Pros: Easy to explain. "We are off by \$5k on average."

Cons: Treats all errors equally.

MSE (Mean Squared Error)

Average of squared errors.

Pros: Penalizes outliers heavily. Large errors become HUGE.

Cons: Harder to interpret (units are squared).

```
from sklearn.metrics import mean_squared_error,  
mean_absolute_error  
print("MAE:", mean_absolute_error(y_test, predictions))  
print("MSE:", mean_squared_error(y_test, predictions))
```

Metrics: RMSE & R-Squared

RMSE

Root Mean Squared Error. Just the square root of MSE. It brings the units back to normal (e.g., dollars), but keeps the penalty for large errors.

R-Squared (R^2)

Score: 0 to 1.

Tells you how much of the variance in the data your model explains. 1.0 is a perfect fit; 0.0 means the model is useless.

| Topic 2: Classification

Predicting Category

Mapping input variables to a **discrete** label.

- Spam vs Ham (Binary)
- Tumor: Benign vs Malignant
- Handwriting: 0-9 (Multi-class)

$$\left\{ \begin{array}{l} P(y=1|x) \\ P(y=0|x) \end{array} \right.$$

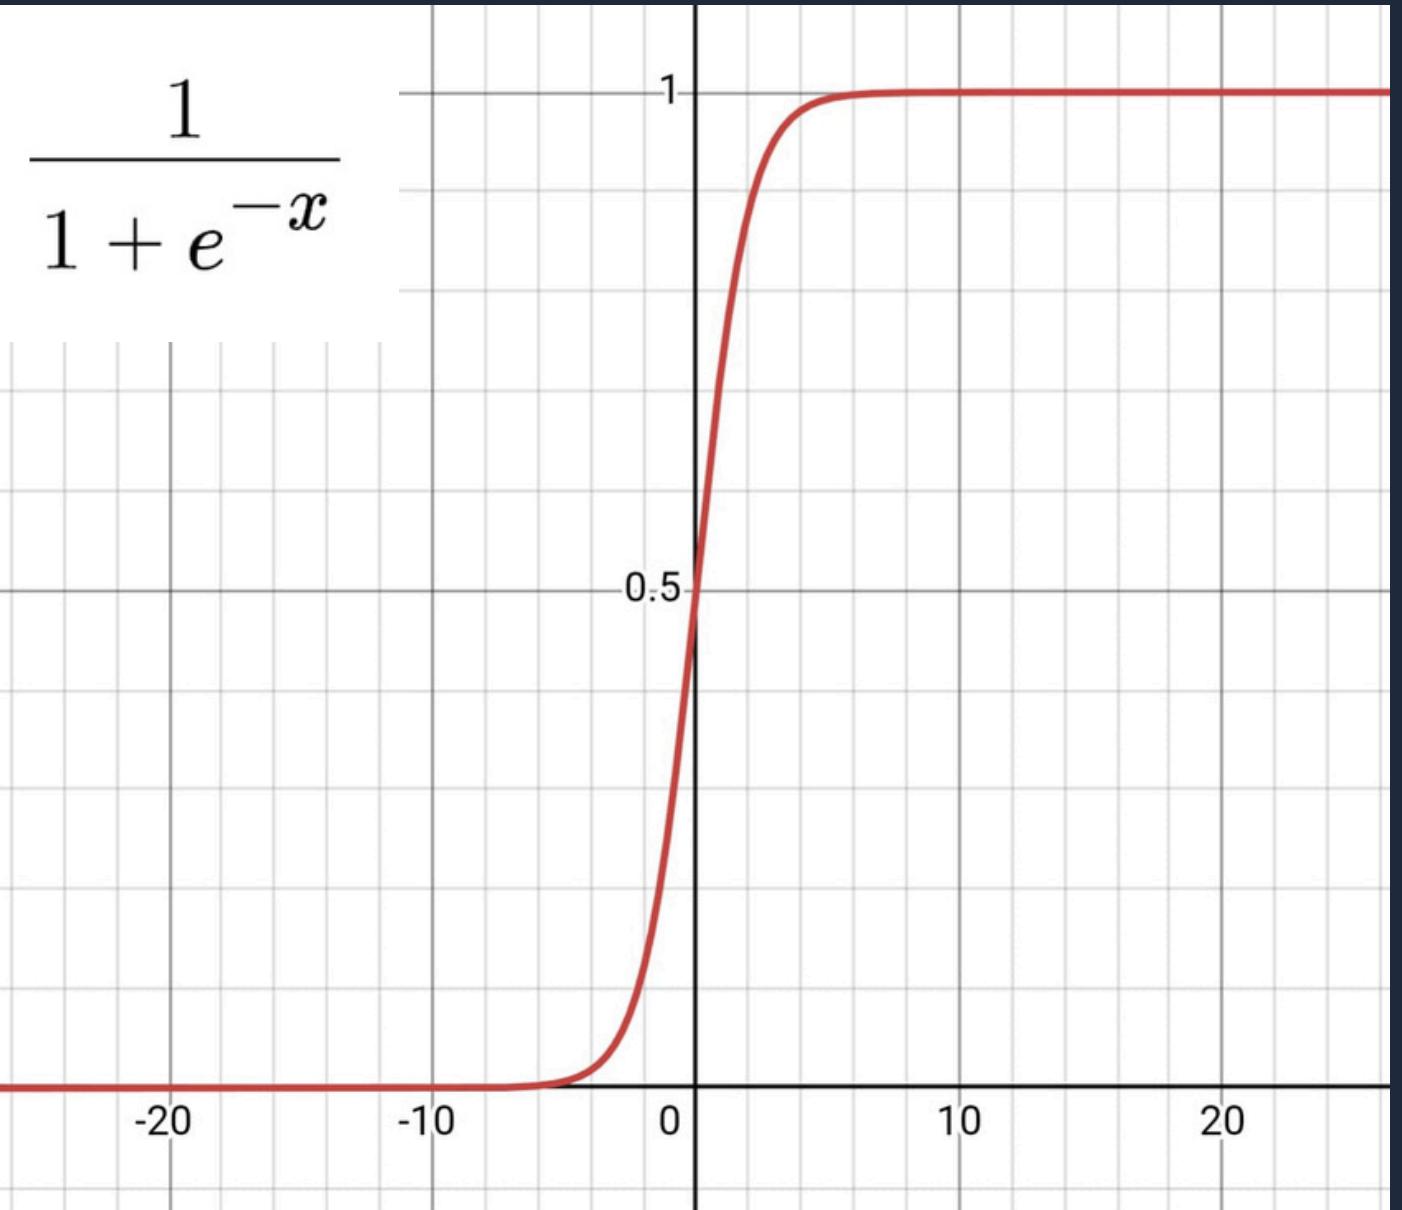
Logistic Regression

It's a Classifier!

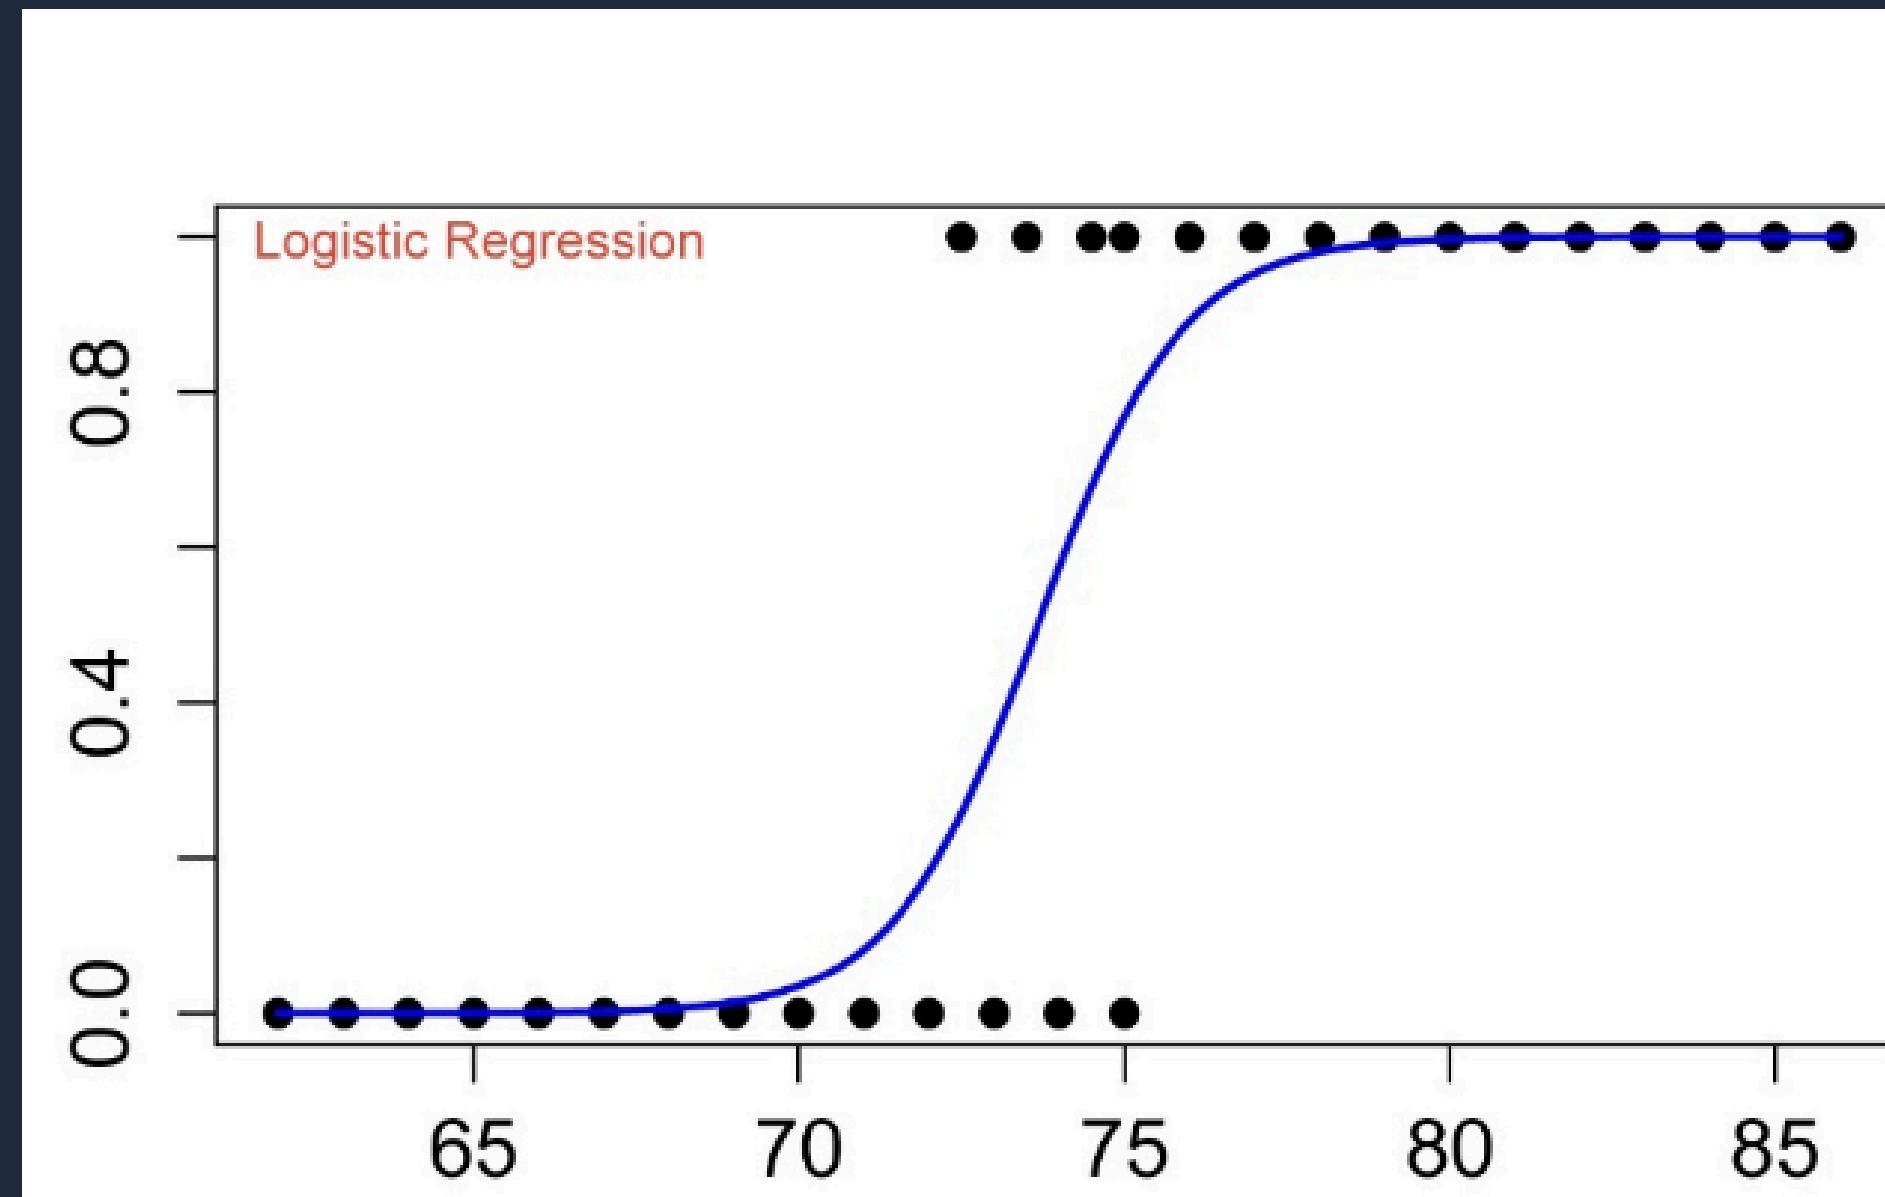
Why not use Linear Regression? Because linear lines go to infinity. Probabilities must stay between 0 and 1.

Solution: Wrap the linear equation in a **Sigmoid Function**.

```
from sklearn.linear_model import LogisticRegression  
clf = LogisticRegression()  
clf.fit(X_train, y_train)  
# Returns 0 or 1 preds = clf.predict(X_test)
```



| The Sigmoid Curve



The "S" curve forces any input number into a safe probability range (0 to 1). We typically set a threshold at 0.5.

K-Nearest Neighbors (KNN)

"Birds of a feather flock together"

KNN assumes similar things exist in close proximity. To classify a new point, it looks at the 'K' nearest points.

Note: Very slow on large datasets because it has to calculate distance to *every* point.

```
from sklearn.neighbors import KNeighborsClassifier  
# K=3 means look at 3 closest neighbors  
knn = KNeighborsClassifier(n_neighbors=3)  
knn.fit(X_train, y_train)
```

| Context: The Impact of 'K'

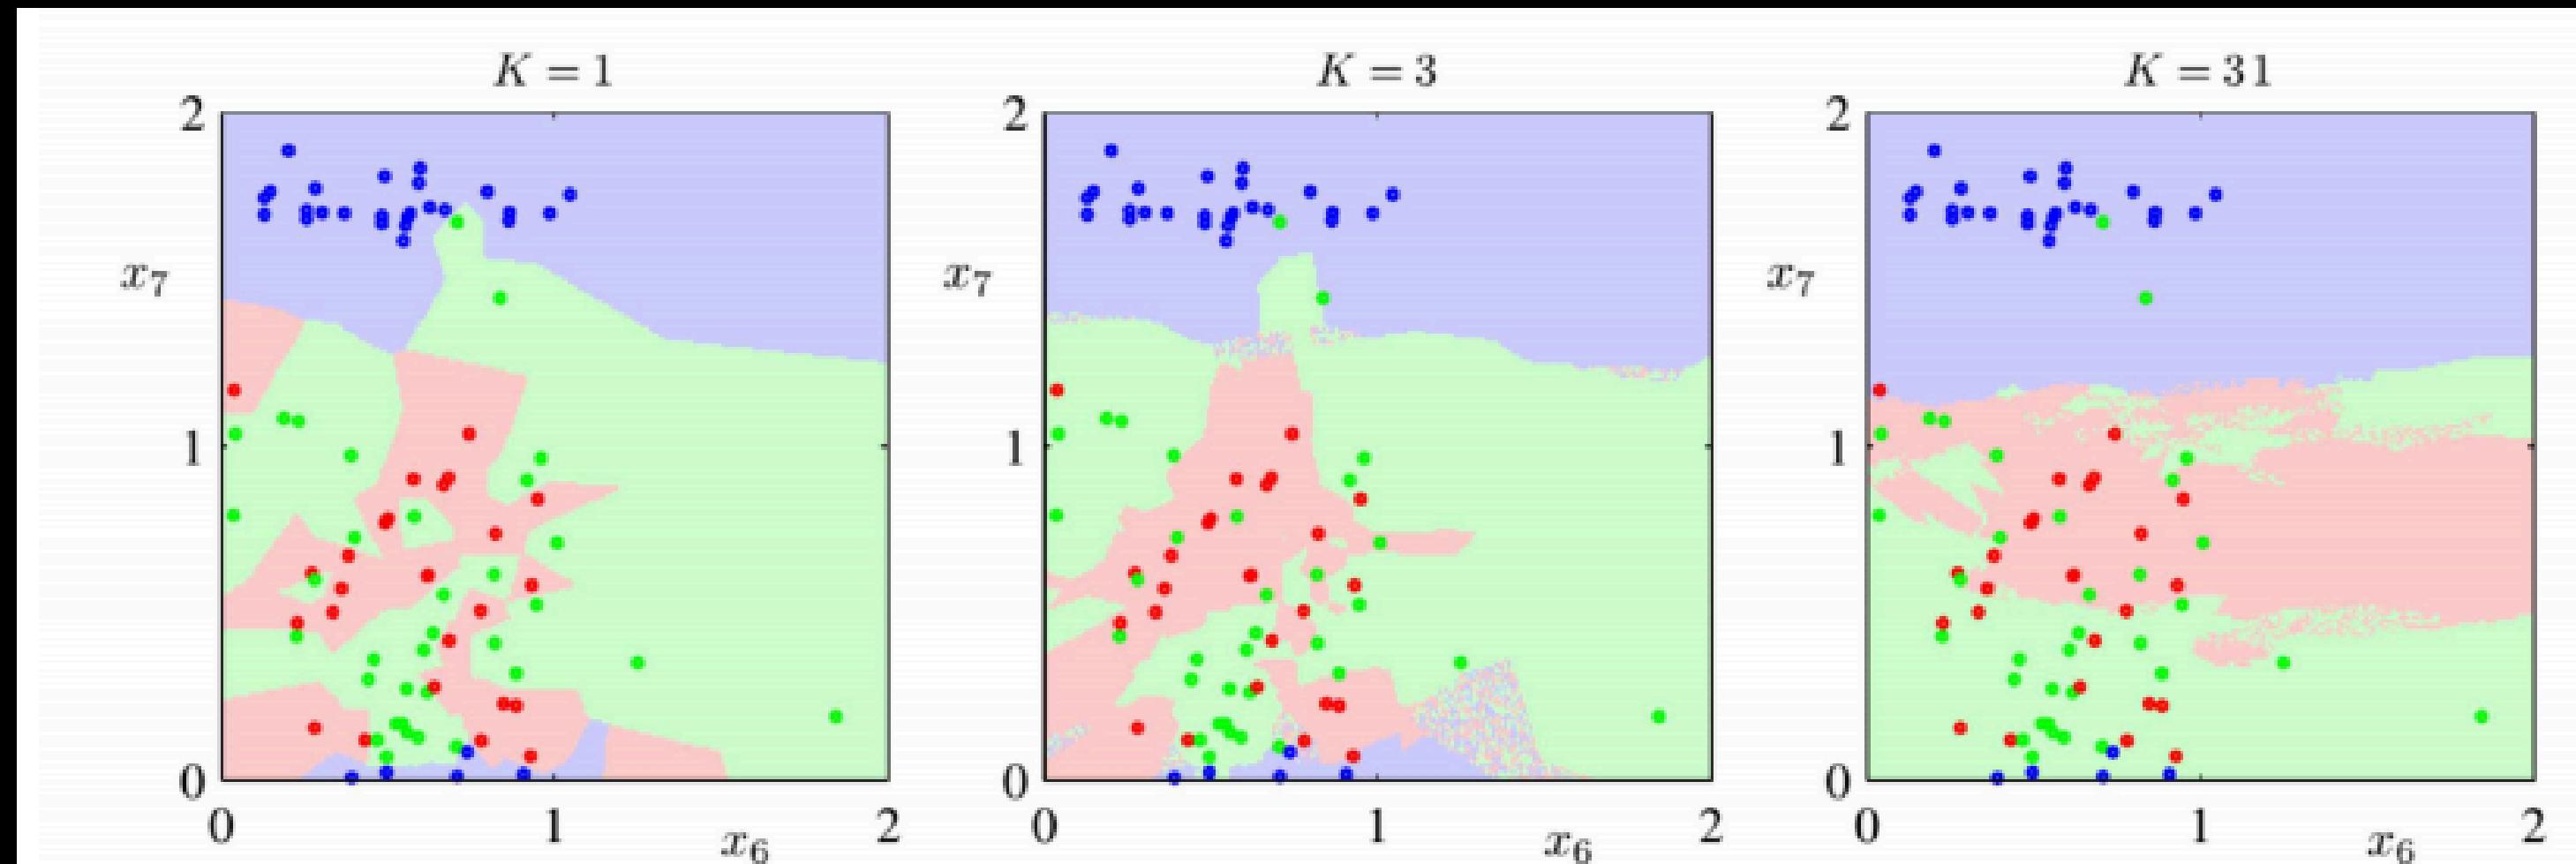
Low K (e.g., K=1)

Overfitting. The boundary is jagged and reacts to every single noise point. If your neighbor is an outlier, you get misclassified.

High K (e.g., K=100)

Underfitting. The boundary becomes too smooth. It ignores local details and just votes for the majority class of the whole area.

KNN Decision Boundaries



Notice how small K creates islands of complex boundaries, while large K creates smooth, simple lines.

Decision Trees

20 Questions

The model asks a series of Yes/No questions to split the data.

Goal: Create "Pure" leaves (groups where everyone belongs to the same class).

```
from sklearn.tree import DecisionTreeClassifier  
# max_depth limits how many questions it can ask  
dt = DecisionTreeClassifier(max_depth=5)  
dt.fit(X_train, y_train)
```

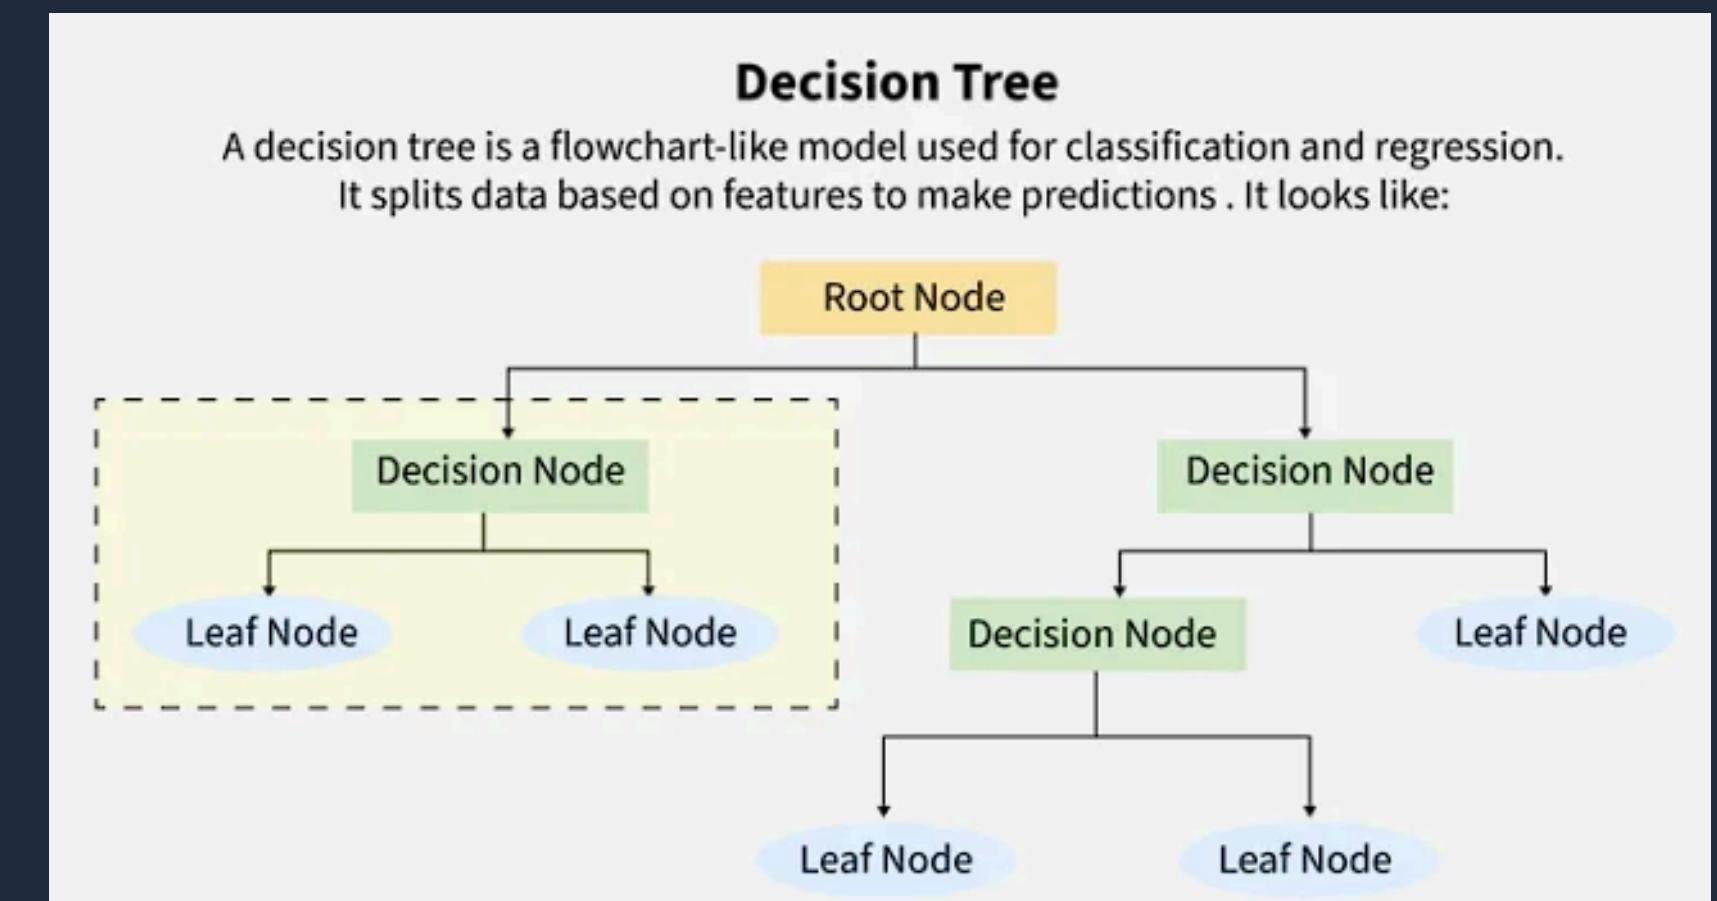
| Context: How it Splits

Entropy / Gini Impurity

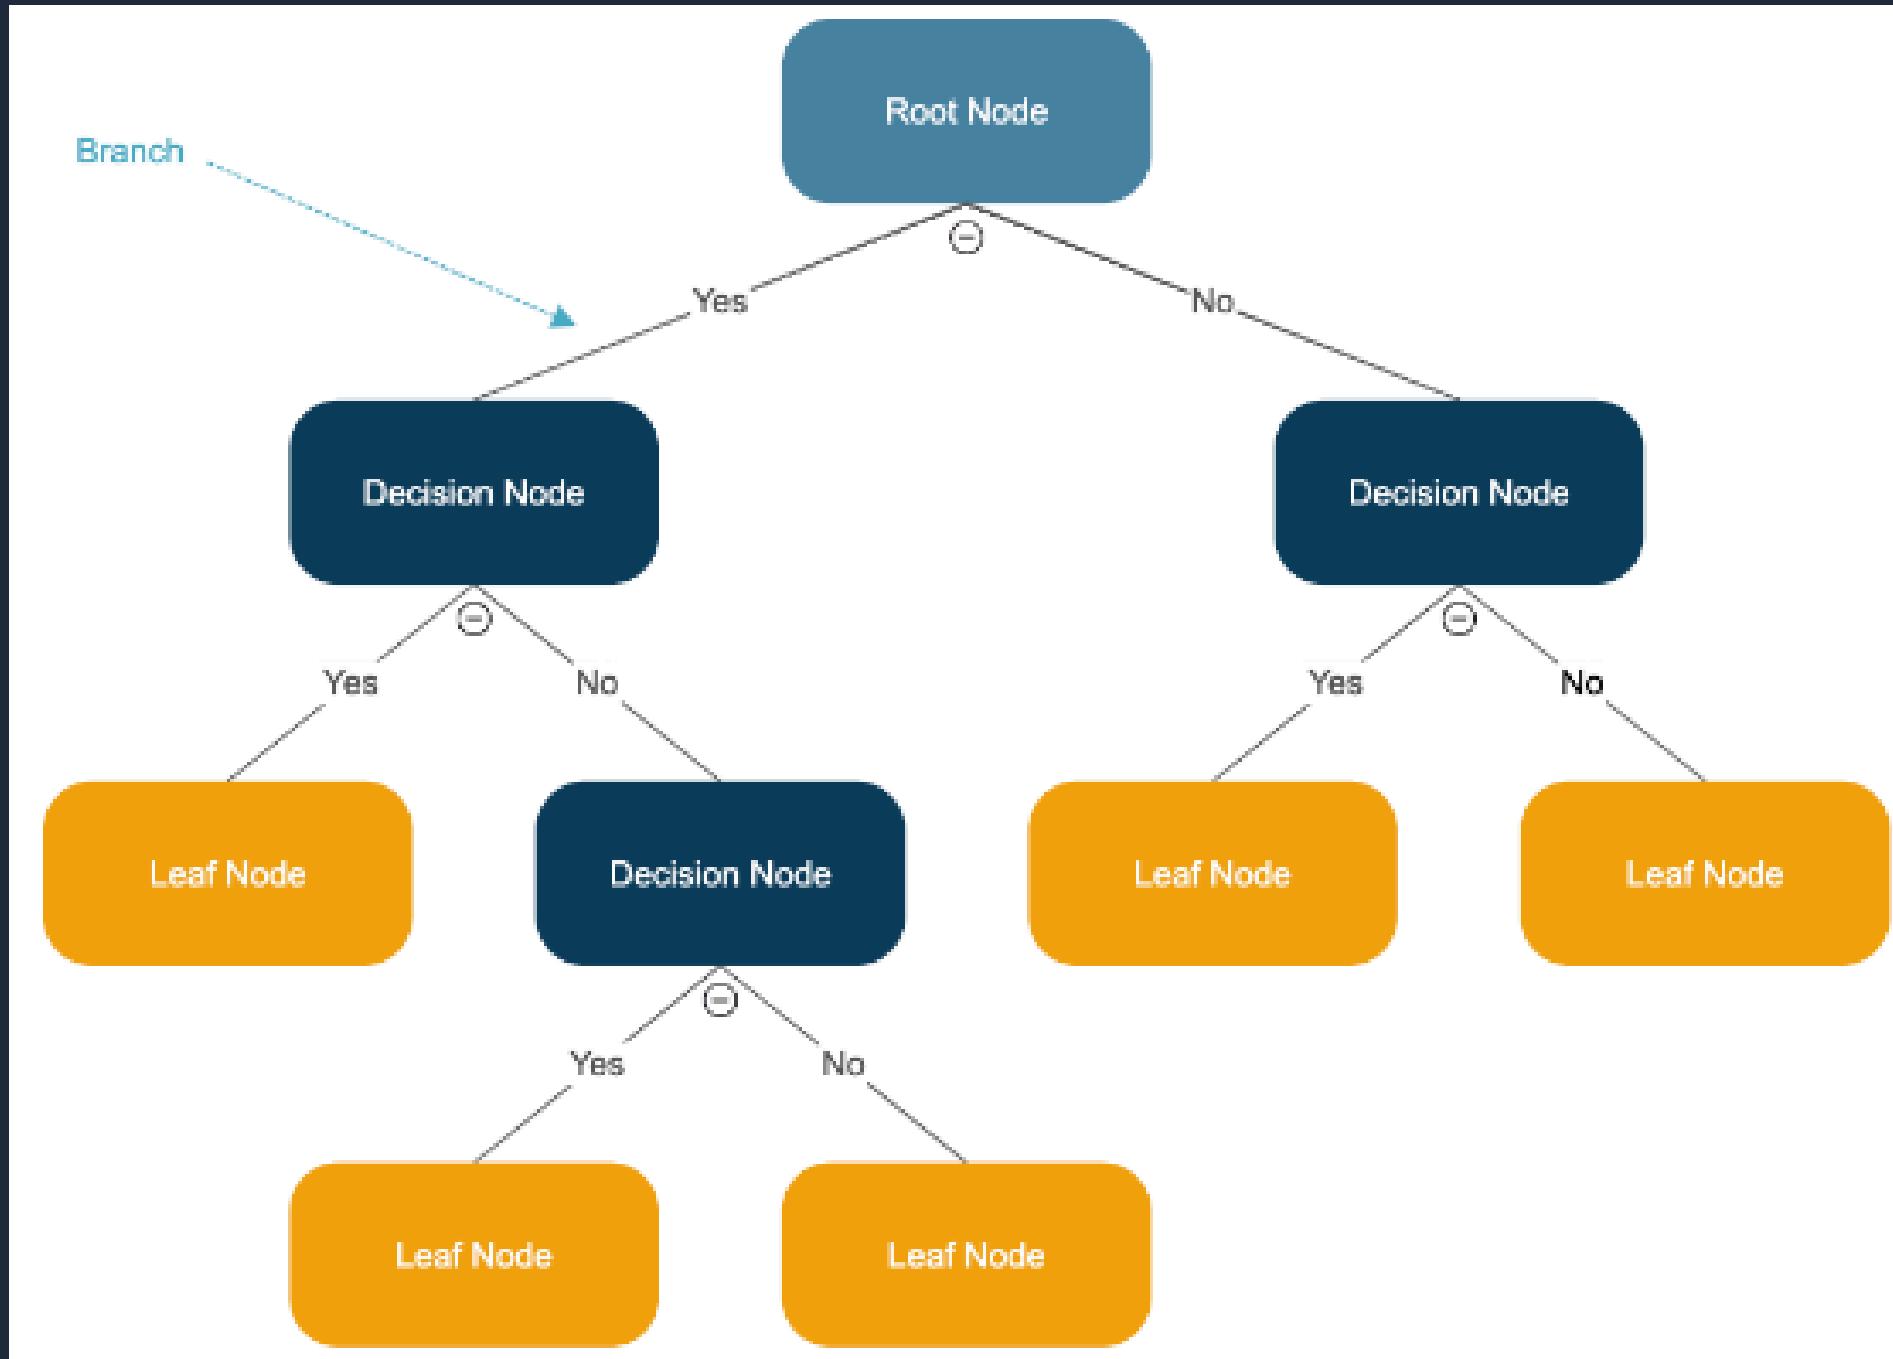
The tree calculates a score to measure "Messiness".

- **High Entropy:** 50% Cats, 50% Dogs (Messy).
- **Low Entropy:** 100% Cats (Pure).

The algorithm greedily chooses the question (e.g., "Is Weight > 10kg?") that reduces entropy the most.



Tree Visualization



Interpretable AI: Unlike neural networks, you can simply follow the path to understand why the model made a decision.

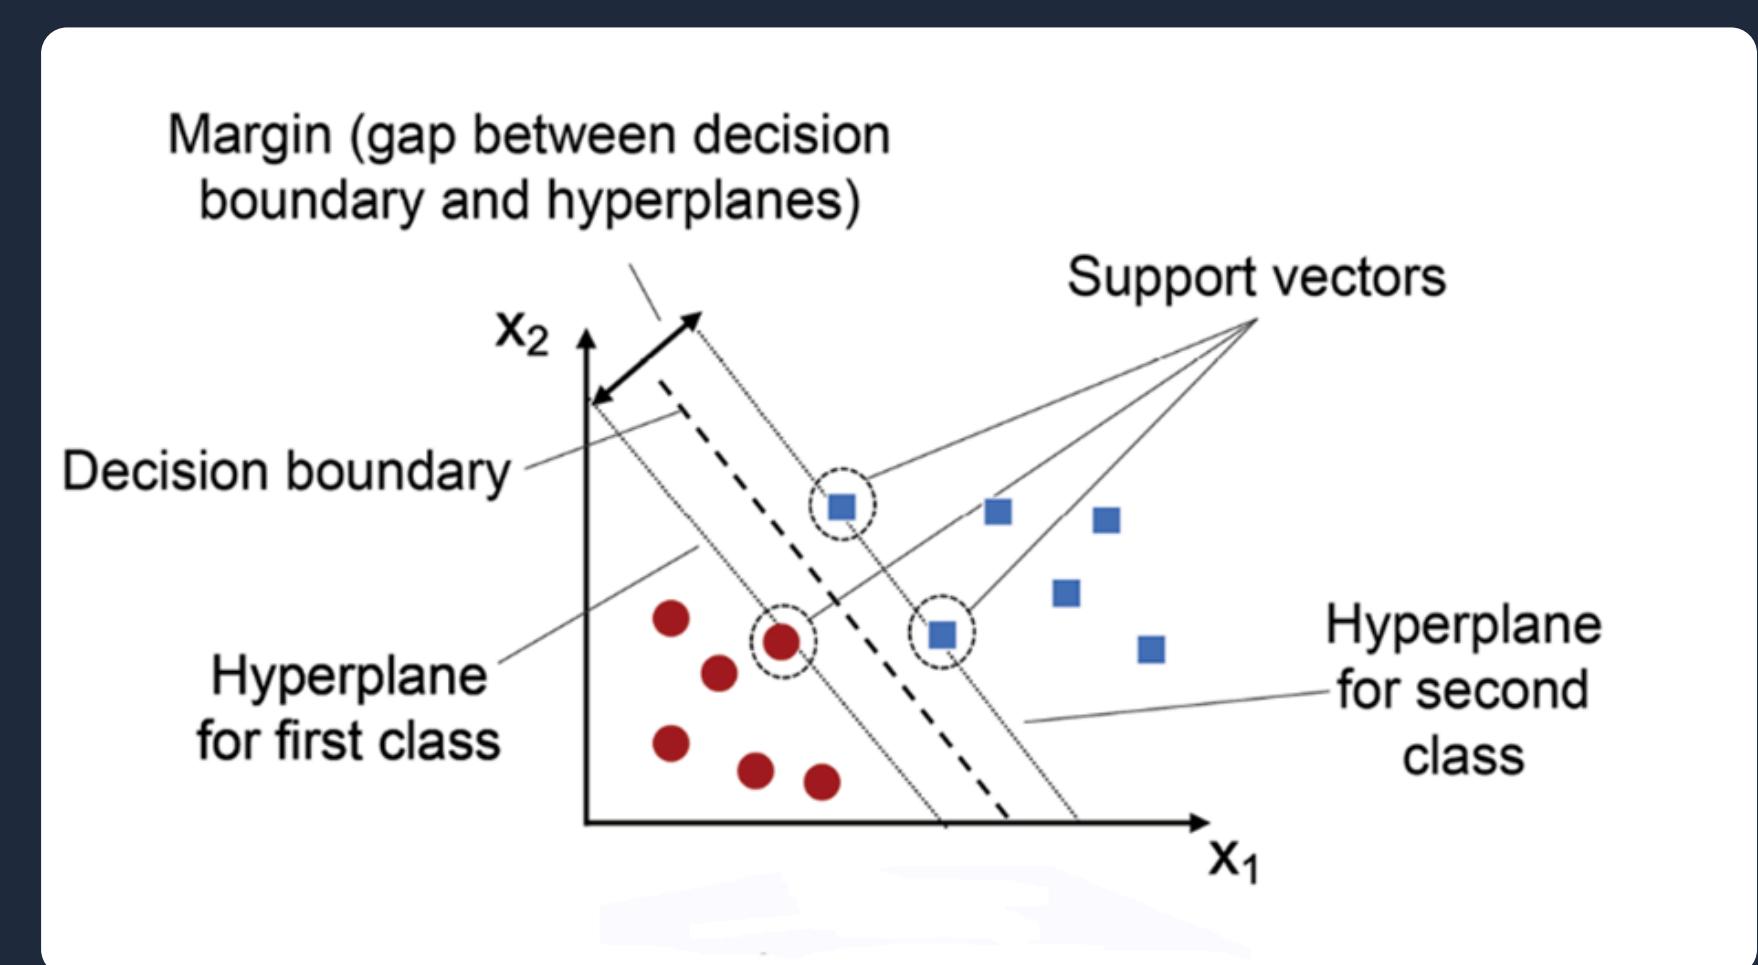
Support Vector Machines (SVM)

The Widest Street

SVM doesn't just find *a* line that separates classes; it finds the *best* line.

Margin: The distance between the boundary and the closest points (Support Vectors). SVM maximizes this margin.

```
from sklearn.svm import SVC  
# kernel='linear' for straight lines  
svm = SVC(kernel='linear')  
svm.fit(X_train, y_train)
```



| Context: The Kernel Trick

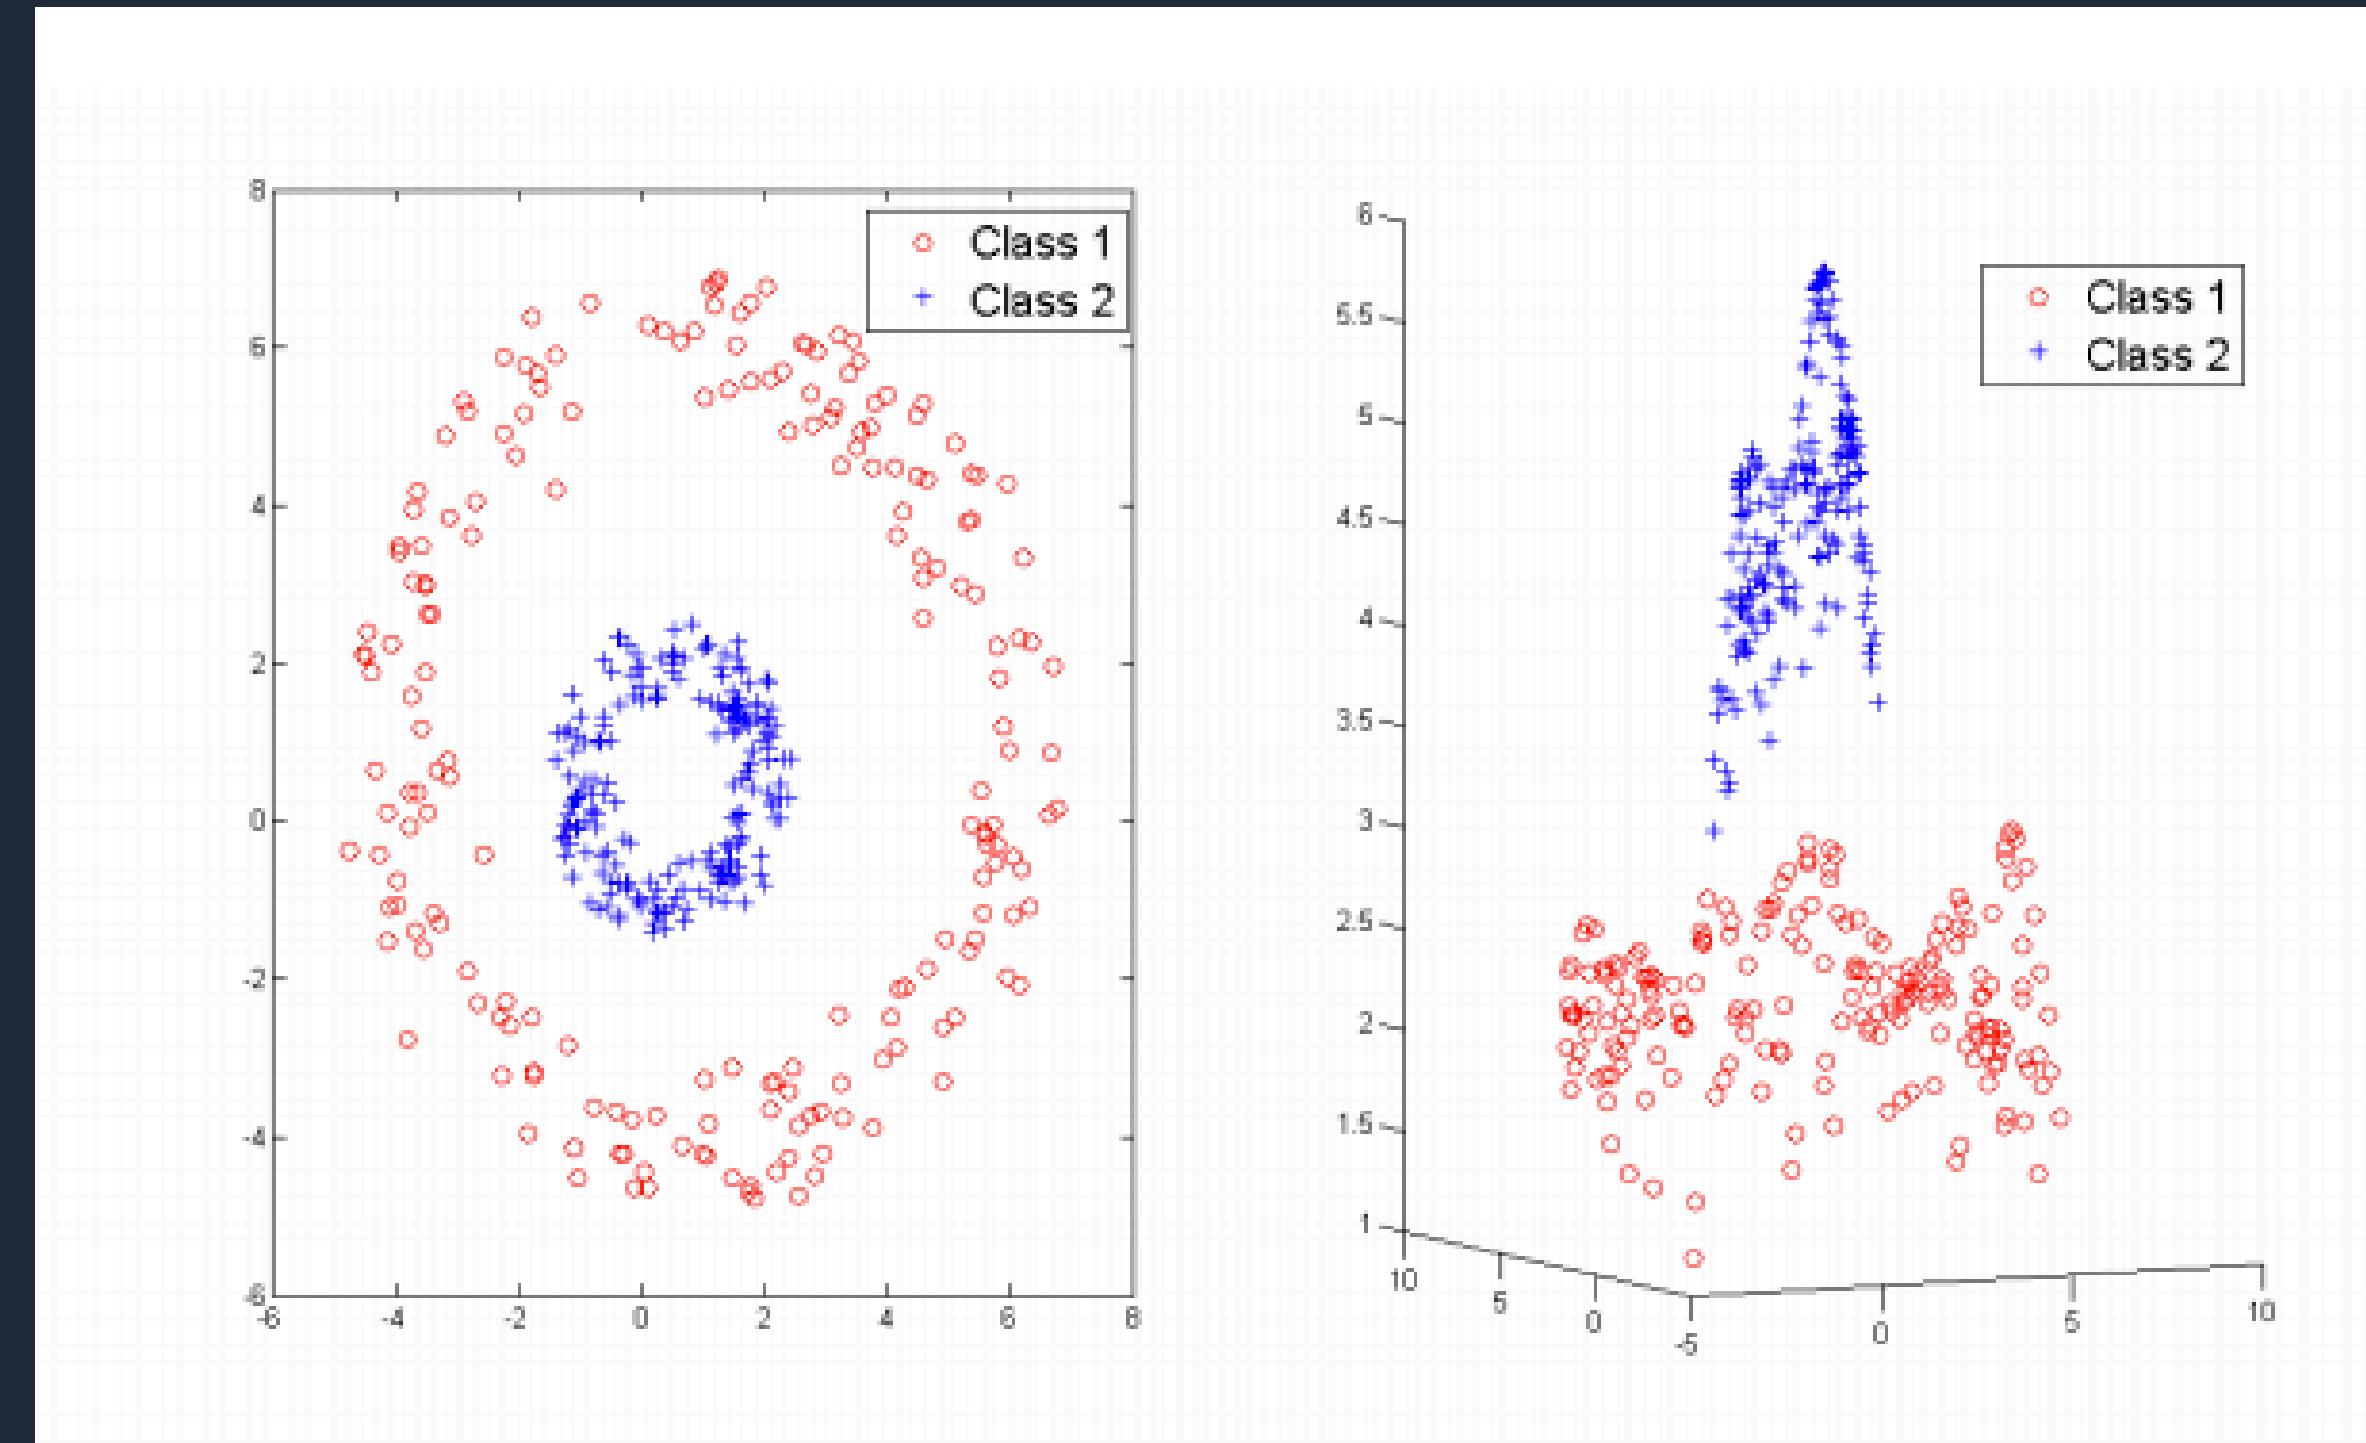
What if data isn't separable?

Imagine red dots in the center and blue dots surrounding them. A straight line can't split them.

Solution: Project data into 3D! In 3D, you might be able to slide a flat sheet (hyperplane) between them. When you project back to 2D, the boundary looks like a circle.

```
# 'rbf' is the default kernel for non-linear data
svm = SVC(kernel='rbf')
```

Visualizing 2D to 3D Projection



By lifting the data into higher dimensions, complex non-linear problems become simple linear ones.

Evaluation: Confusion Matrix

The 4 Outcomes

Accuracy hides details. We need to know *how* we are wrong.

- **True Positive (TP)**: Correctly spotted Fire.
- **True Negative (TN)**: Correctly ignored Safe.
- **False Positive (FP)**: False Alarm (Type I).
- **False Negative (FN)**: Missed Fire (Type II - Dangerous!).

Confusion Matrix

| | Actually Positive (1) | Actually Negative (0) |
|------------------------|-----------------------|-----------------------|
| Predicted Positive (1) | True Positives (TPs) | False Positives (FPs) |
| Predicted Negative (0) | False Negatives (FNs) | True Negatives (TNs) |

| Context: Precision vs Recall

Precision

"Of all the alarms we rang, how many were real fires?"

Optimize this to avoid spamming users (False Positives).

$$P = \frac{TP}{(TP + FP)}$$

Recall

"Of all the real fires, how many did we detect?"

Optimize this for cancer detection (False Negatives are fatal).

$$R = \frac{TP}{(TP + FN)}$$

The F1-Score

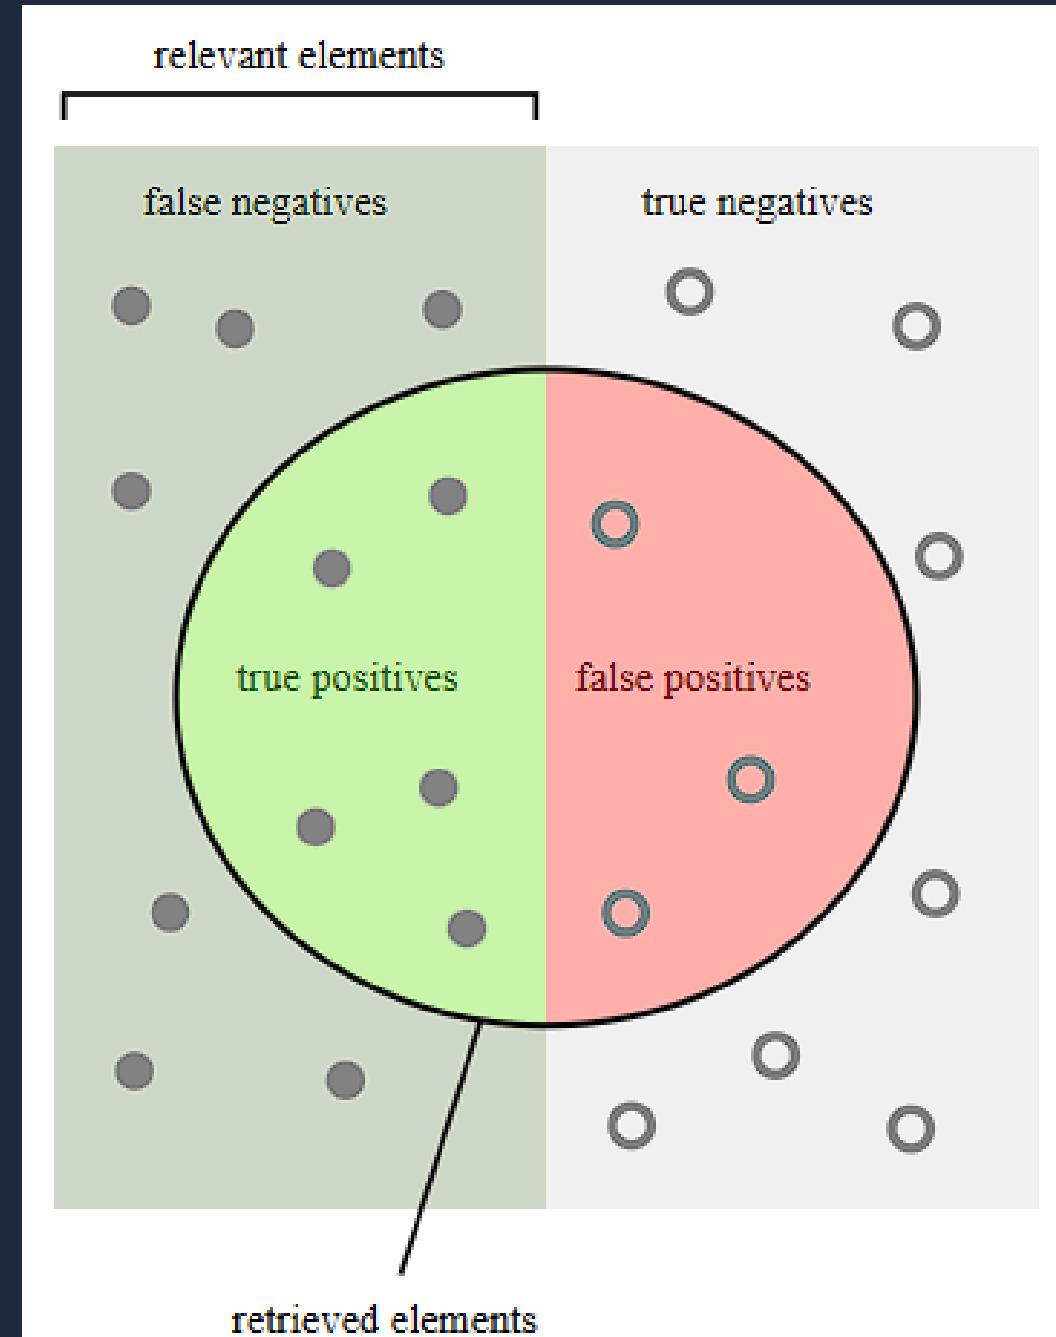
Balancing the Two

You can't have perfect Precision and Recall simultaneously.

As you catch more cases (Recall \uparrow), you inevitably make more false alarms (Precision \downarrow).

F1-Score: The Harmonic Mean of Precision and Recall. Use this metric when classes are imbalanced.

$$F1 = \frac{2^*P^*R}{(P+R)}$$



How many retrieved items are relevant?

$$\text{Precision} = \frac{\text{true positives}}{\text{true positives} + \text{false positives}}$$

How many relevant items are retrieved?

$$\text{Recall} = \frac{\text{true positives}}{\text{true positives} + \text{false negatives}}$$

Evaluation Code

```
from sklearn.metrics import classification_report, confusion_matrix  
print(confusion_matrix(y_test, preds)) # Prints Precision, Recall, F1 for each class  
print(classification_report(y_test, preds))
```

Masterclass Summary

Regression

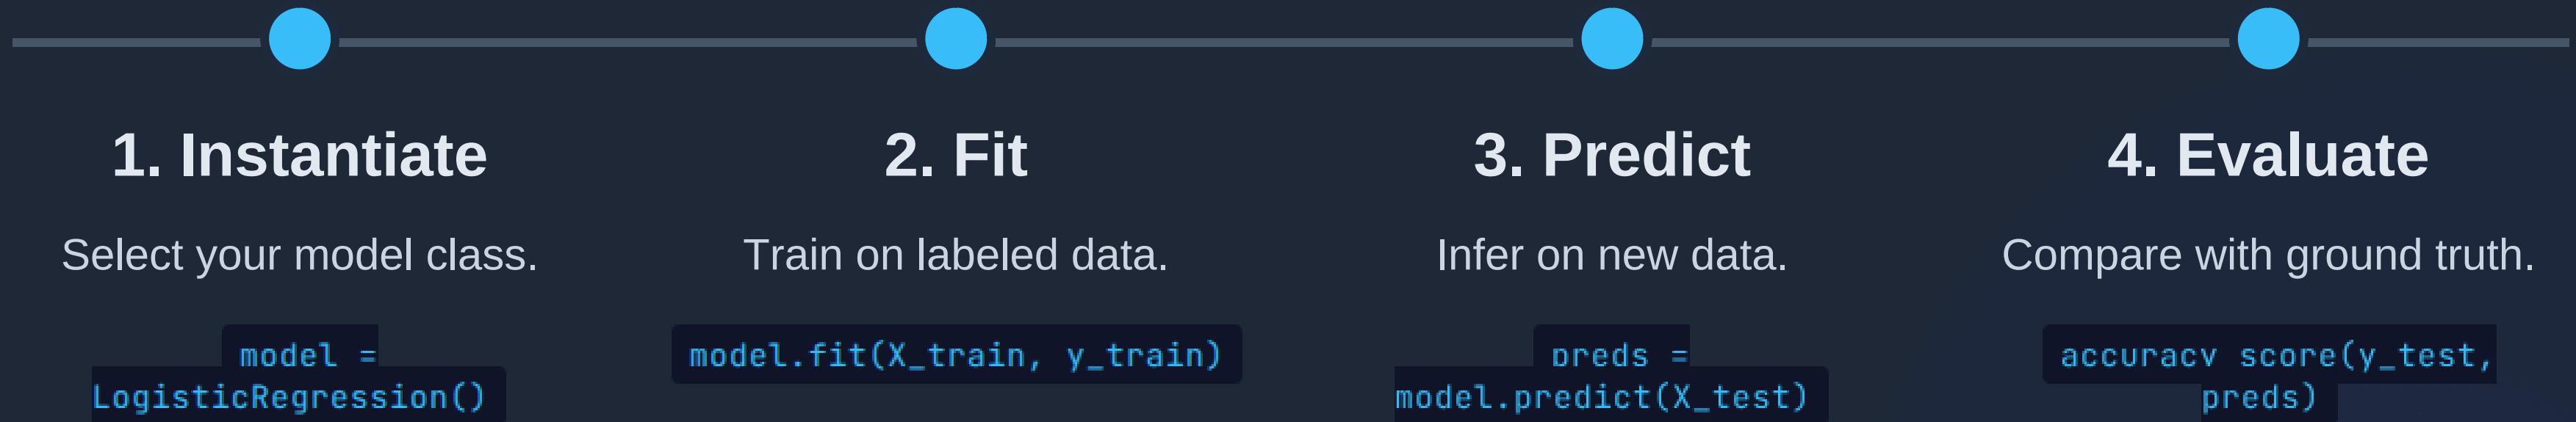
- Goal: Continuous #
- Alg: Linear, Polynomial
- Metric: MSE, R2
- Key: Gradient Descent

Classification

- Goal: Category Label
- Alg: LogReg, KNN, SVM, Tree
- Metric: F1, Confusion Matrix
- Key: Decision Boundary

Standard Library Workflow

Whether using Scikit-Learn, PyTorch, or TensorFlow, the pattern remains consistent.



Questions?

Thank you for attending.